

HW 1

Shivalika Chavan

2026-02-15

```
housing_train = read.csv("./data/housing_training.csv") |>
  janitor::clean_names() |>
  mutate(
    overall_qual = factor(overall_qual,
      levels = c("Very_Excellent",
                 "Excellent",
                 "Very_Good",
                 "Good",
                 "Above_Average",
                 "Average",
                 "Below_Average",
                 "Fair"),
      ordered = TRUE),
    kitchen_qual = factor(kitchen_qual,
      levels = c("Excellent",
                 "Good",
                 "Typical",
                 "Fair"),
      ordered = TRUE),
    fireplace_qu = factor(fireplace_qu,
      levels = c("Excellent",
                 "Good",
                 "Fair",
                 "Poor",
                 "No_Fireplace"
      ),
      ordered = TRUE),
    fireplace_qu = replace_na(fireplace_qu, "No_Fireplace"),
    exter_qual = factor(exter_qual,
      levels = c("Excellent",
                 "Good",
                 "Typical",
                 "Fair"),
      ordered = TRUE)
  )

housing_test = read.csv("./data/housing_test.csv") |>
  janitor::clean_names() |>
  mutate(
    overall_qual = factor(overall_qual,
      levels = c("Very_Excellent",
                 "Excellent",
                 "Very_Good",
                 "Good",
                 "Above_Average",
                 "Average",
                 "Below_Average",
                 "Fair"),
      ordered = TRUE))
```

```

        "Very_Good",
        "Good",
        "Above_Average",
        "Average",
        "Below_Average",
        "Fair"),
    ordered = TRUE),
kitchen_qual = factor(kitchen_qual,
    levels = c("Excellent",
              "Good",
              "Typical",
              "Fair"),
    ordered = TRUE),
fireplace_qu = factor(fireplace_qu,
    levels = c("Excellent",
              "Good",
              "Fair",
              "Poor",
              "No_Fireplace"
            ),
    ordered = TRUE),
fireplace_qu = replace_na(fireplace_qu, "No_Fireplace"),
exter_qual = factor(exter_qual,
    levels = c("Excellent",
              "Good",
              "Typical",
              "Fair"),
    ordered = TRUE)
)

y = housing_train |> pull(sale_price)
x = model.matrix(sale_price ~ ., housing_train)[,-1]

y_test = housing_test$sale_price
x_test = model.matrix(sale_price ~ ., housing_test)[,-1]

```

a - LASSO Model

```

set.seed(1234)

lambda = 10^(seq(-5, 5, 0.2)) # lambda grid of lambda values for penalty tuning

# k-fold cross-validation for Lasso (alpha = 1) over the lambda values
lasso_cv = cv.glmnet(x, y,
                      alpha = 1,
                      lambda = lambda)

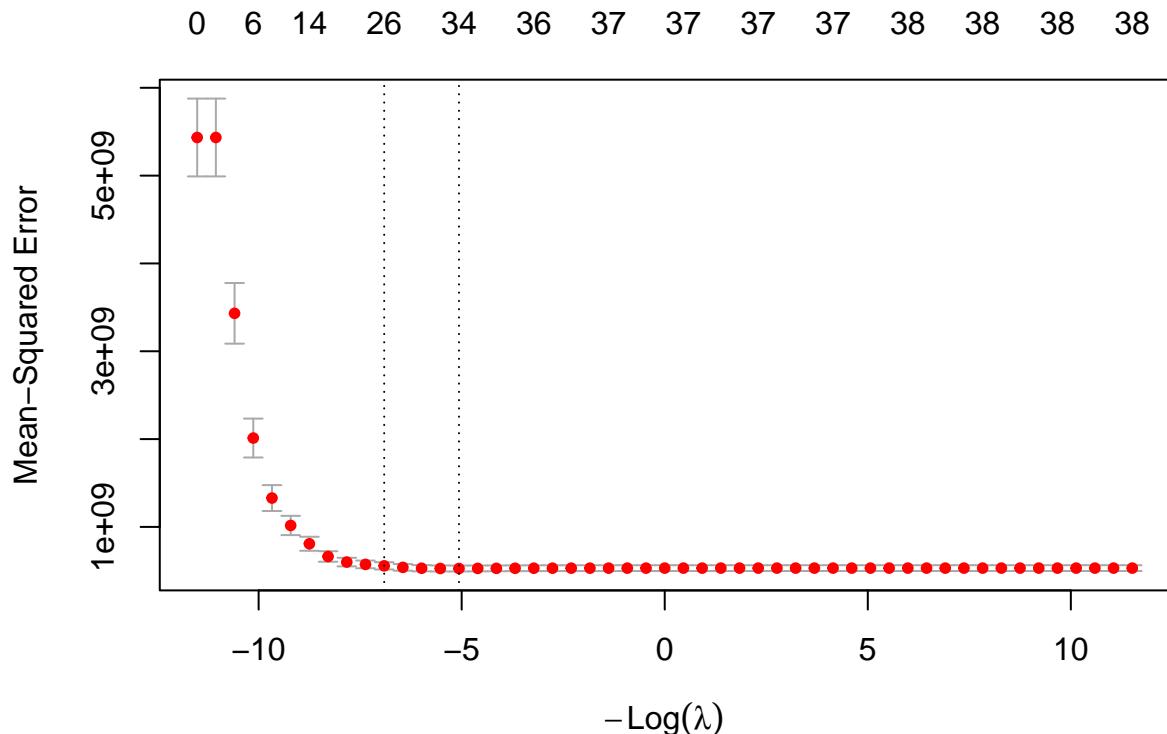
lambda_min_glmnet = lasso_cv[["lambda.min"]] # minimum mean cross-validated error
lambda_1se_glmnet = lasso_cv[["lambda.1se"]] # largest value of lambda such that error is within 1 stan
CVM_min = lasso_cv |> broom::tidy() |> filter(lambda == lambda_min_glmnet) |> pull(estimate)
CVM_1se = lasso_cv |> broom::tidy() |> filter(lambda == lambda_1se_glmnet) |> pull(estimate)

```

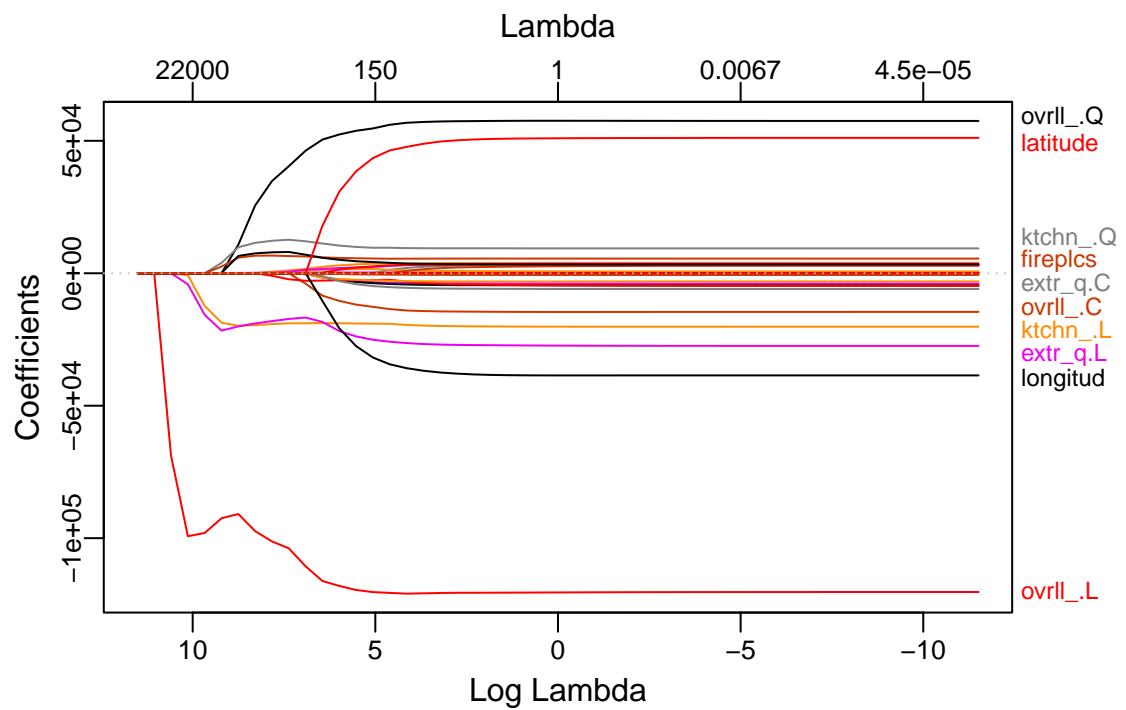
The λ value with the smallest CVM (5.2645935×10^8) is 158.4893192. The λ value with the 1SE CVM

(5.5664618×10^8) is 1000.

```
plot(lasso_cv)
```



```
plot_glmnet(lasso_cv$glmnet.fit)
```



Test Error with $\lambda = 158.49$:

```
y_pred = predict(lasso_cv, newx = x_test, s = lambda_min_glmnet, type = "response")
mspe_lasso <- mean((y_test - y_pred)^2)
```

The test error is 4.3060946×10^8 .

```
coef_1se = predict(lasso_cv, type = "coefficients", s = lambda_1se_glmnet)
# Count non-zero coefficients to determine the number of predictors (excluding intercept)
num_predictors_1se = sum(coef_1se != 0) - 1
```

When using λ_{1SE} , there are 26 predictors.

b - Elastic Net Model

```
set.seed(1234)

alpha = seq(0, 1, length = 21) # alpha grid ranging from 0 (Ridge) to 1 (Lasso)
lambda = 10^(seq(-5, 5, 0.2)) # lambda grid of lambda values for penalty tuning

# Iterate through each alpha to perform cross-validation and store results in a tibble
enet_cv_results = tibble(alpha = alpha) |>
  mutate(
    cv_fit = map(alpha, ~cv.glmnet(x, y, alpha = .x, lambda = lambda)), # runs glmnet for each alpha
    min_cvm = map_dbl(cv_fit, ~min(.x$cvm)) # finds min CVM for each model at each alpha
  )

# pull the alpha value with the lowest overall CVM
enet_alpha_min_cvm = enet_cv_results |>
  filter(min_cvm == min(min_cvm)) |>
  pull(alpha)

# pull the corresponding model with that optimal alpha
best_enet_cv_fit = enet_cv_results |>
  filter(alpha == enet_alpha_min_cvm) |>
  pull(cv_fit) |>
  pluck(1)

# pull 1SE lambda for optimal alpha
lambda_1se_enet = best_enet_cv_fit$lambda.1se

# make predictions using optimal alpha and 1SE lambda at that alpha
y_pred = predict(best_enet_cv_fit, newx = x_test, s = "lambda.1se")
mspe_elastic_net = mean((y_test - y_pred)^2)
```

The selected tuning parameters are $\alpha = 0.75$ and $\lambda_{1SE} = 1000$. The model with these parameters has a test error of 4.2461127×10^8 .

Applying the 1SE rule is a bit more complicated with elastic net, because it has two parameters, α and λ . The 1SE rule, when used in the Lasso model, is easy to implement because a larger value for λ means fewer non-zero coefficients and a more parsimonious model. In elastic net models, this will work when α is fixed at a single value.

c - Partial Least Squares (PLS) Model

```

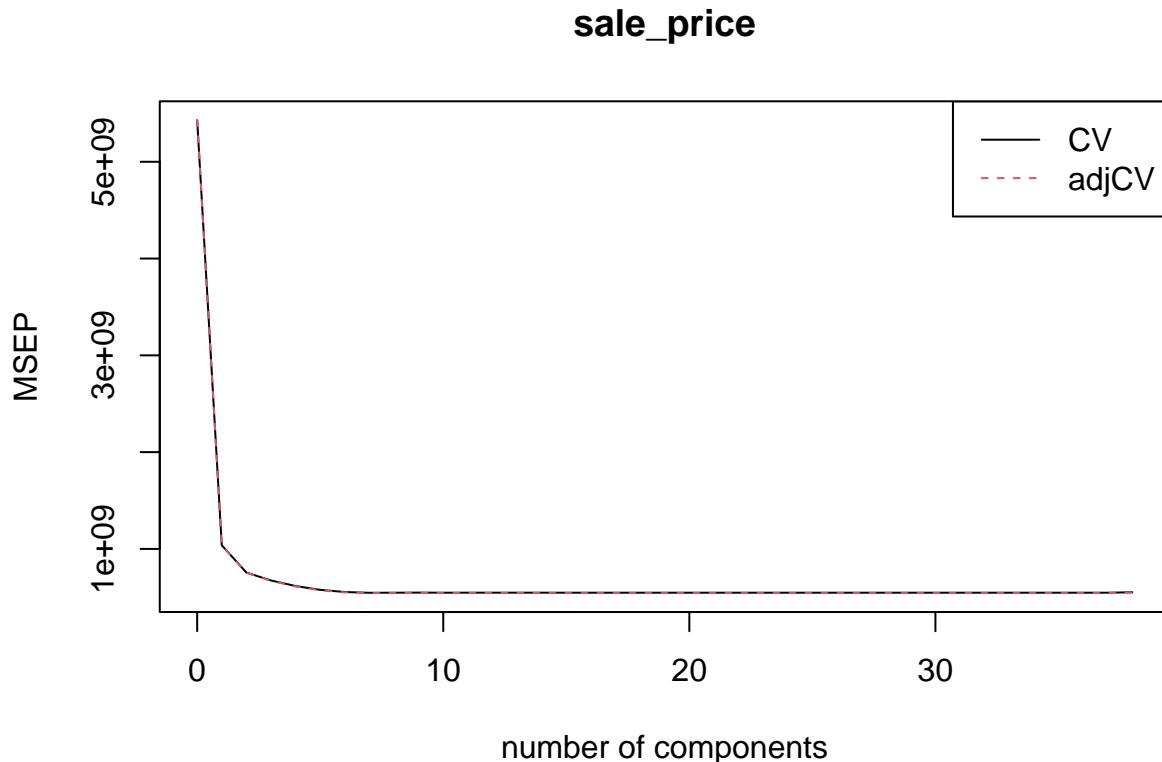
set.seed(1234)

pls_mod <- plsr(sale_price ~ .,
                  data = housing_train,
                  scale = TRUE, # similar scaling importance as PCR
                  validation = "CV")
summary(pls_mod)

## Data: X dimension: 1440 38
## Y dimension: 1440 1
## Fit method: kernelpls
## Number of components considered: 38
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##          (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV         73685    32214   27493   25962   24834   24019   23548
## adjCV      73685    32199   27454   25886   24730   23928   23476
##          7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV         23381    23400   23419   23395   23401   23403   23404
## adjCV      23311    23326   23342   23319   23324   23325   23326
##          14 comps 15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## CV         23406    23398   23399   23399   23400   23399   23399
## adjCV      23327    23320   23321   23321   23322   23321   23321
##          21 comps 22 comps 23 comps 24 comps 25 comps 26 comps 27 comps
## CV         23399    23399   23399   23399   23399   23399   23399
## adjCV      23321    23321   23321   23321   23321   23321   23321
##          28 comps 29 comps 30 comps 31 comps 32 comps 33 comps 34 comps
## CV         23399    23399   23399   23399   23399   23399   23399
## adjCV      23321    23321   23321   23321   23321   23321   23321
##          35 comps 36 comps 37 comps 38 comps
## CV         23399    23399   23399   23499
## adjCV      23321    23321   23321   23400
##
## TRAINING: % variance explained
##          1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps
## X          22.75   28.96   34.05   38.43   42.42   46.49   49.60
## sale_price 81.33   86.89   88.79   90.02   90.65   90.90   91.02
##          8 comps 9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X          52.36   55.49   57.47   59.89   62.33   64.09   65.6
## sale_price 91.05   91.07   91.09   91.09   91.10   91.10   91.1
##          15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## X          66.94   68.8    70.45   71.54   73.23   74.72
## sale_price 91.10   91.1    91.10   91.10   91.10   91.10
##          21 comps 22 comps 23 comps 24 comps 25 comps 26 comps
## X          76.57   77.8    79.91   81.9    83.09   84.17
## sale_price 91.10   91.1    91.10   91.1    91.10   91.10
##          27 comps 28 comps 29 comps 30 comps 31 comps 32 comps
## X          85.82   87.25   88.32   89.51   90.68   91.96
## sale_price 91.10   91.10   91.10   91.10   91.10   91.10
##          33 comps 34 comps 35 comps 36 comps 37 comps 38 comps
## X          93.66   95.02   97.15   97.97   100.0   100.73
## sale_price 91.10   91.10   91.10   91.10   91.1    91.08

```

```
# plot cross-validated MSEP for PLS
validationplot(pls_mod, val.type = "MSEP", legendpos = "topright")
```



```
# determine the optimal number of components
cv_mse <- RMSEP(pls_mod)
ncomp_cv <- which.min(cv_mse$val[1, ,]) - 1
ncomp_cv
```

```
## 7 comps
##    7
# calculate test MSE
predy2_pls <- predict(pls_mod, newdata = housing_test,
                      ncomp = ncomp_cv)

mspe_pls = mean((y_test - predy2_pls)^2)
```

There are 7 components in the partial least squares model, with an MSPE of 4.5063981×10^8 .

d - Comparing Models

```
summary = tibble(
  model = c("LASSO", "Elastic Net", "PLS"),
  mspe = c(mspe_lasso, mspe_elastic_net, mspe_pls)
)

knitr::kable(summary)
```

model	mspe
LASSO	430609458
Elastic Net	424611266
PLS	450639809

Comparing the mean squared predicted error across the three models, the elastic net model is the best model for making predictions. This means that the model likely benefits from the balance of LASSO and Ridge (λ , α) penalties, as compared to the LASSO model, which only has the λ penalty.

e - Retraining LASSO model using caret instead of glmnet

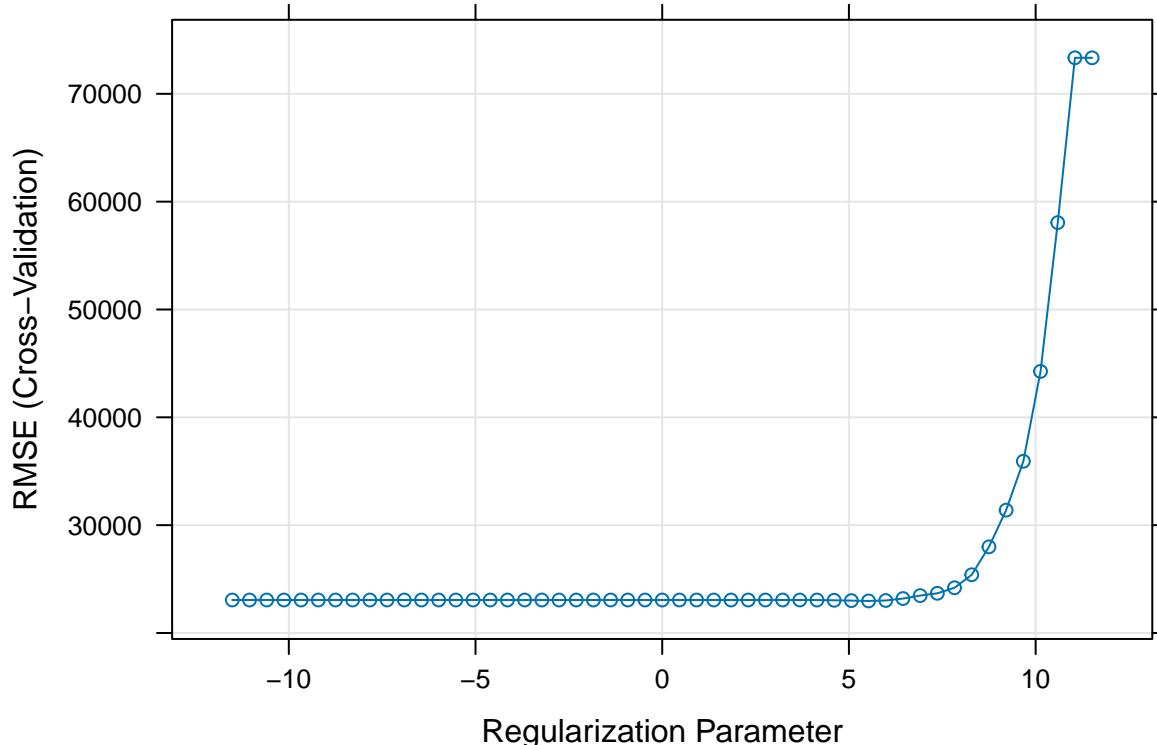
```
set.seed(1234)

ctrl1 = trainControl(method = "cv", number = 10) #cross validation, in this case 10-fold

lasso_caret = train(sale_price ~ .,
                    data = housing_train,
                    method = "glmnet",
                    tuneGrid = expand.grid(alpha = 1,
                                           lambda = lambda),
                    trControl = ctrl1)

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.

plot(lasso_caret, xTrans = log)
```



```
lambda_min_caret = lasso_caret$bestTune$lambda
```

```
tibble(
  package = c("caret (min)", "glmnet (min)", "glmnet (1SE)"),
  lambda = c(lambda_min_caret, lambda_min_glmnet, lambda_1se_glmnet)
) |> knitr::kable(digits = 2)
```

package	lambda
caret (min)	251.19
glmnet (min)	158.49
glmnet (1SE)	1000.00

The “best” values for λ differ slightly between `caret` and `cv.glmnet`. However, these two values are much closer to each other than they are to the 1SE value found using `cv.glmnet`. While they will result in slightly different models, their cross-validation errors will be within 1SE of each other. The reason for this discrepancy is the difference in internal functions; specifically, how each package assigns observations to k cross-validation folds, even when the random seed `set.seed(1234)` is kept consistent.