

HW 1

Shivalika Chavan

2026-02-15

```
housing_train = read.csv("./data/housing_training.csv") |> janitor::clean_names()
housing_test = read.csv("./data/housing_test.csv") |> janitor::clean_names()

y = housing_train |> pull(sale_price)
x = model.matrix(sale_price ~ ., housing_train) [,-1]

x_test = model.matrix(sale_price ~ ., housing_test) [,-1]
y_test = housing_test$sale_price
```

a - LASSO Model

```
set.seed(1234)

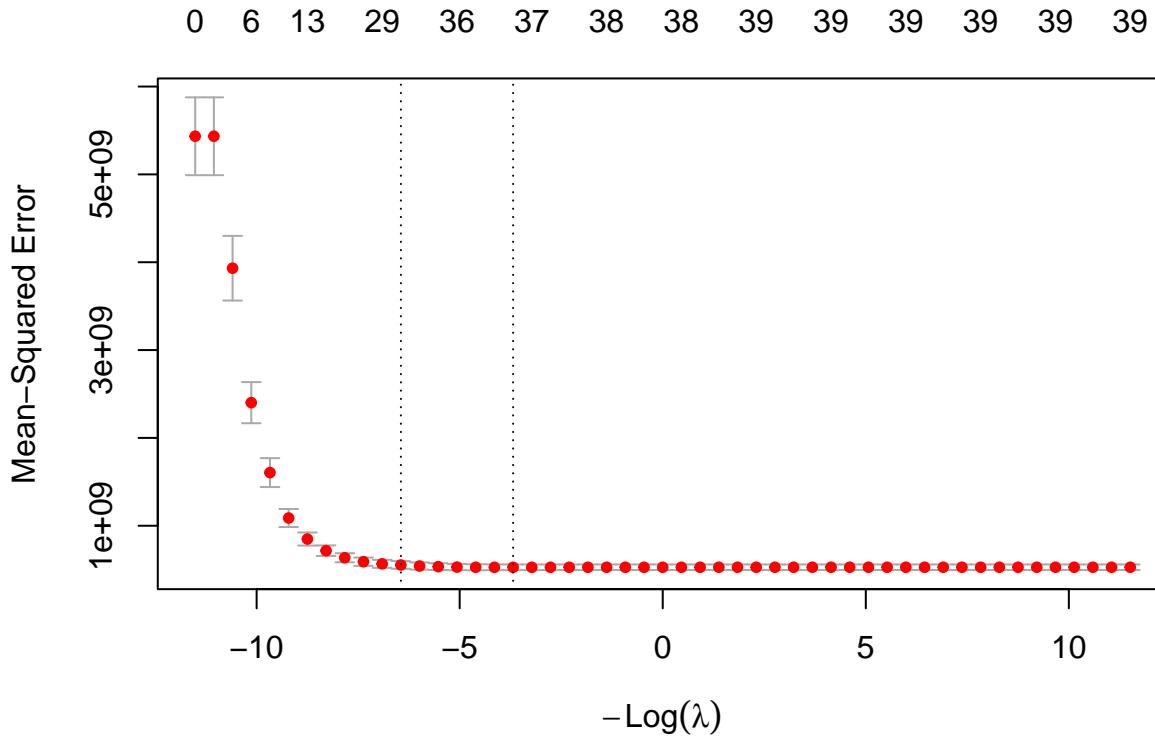
lambda = 10^(seq(-5, 5, 0.2)) # lambda grid of lambda values for penalty tuning

# k-fold cross-validation for Lasso (alpha = 1) over the lambda values
lasso_cv = cv.glmnet(x, y,
                      alpha = 1,
                      lambda = lambda)

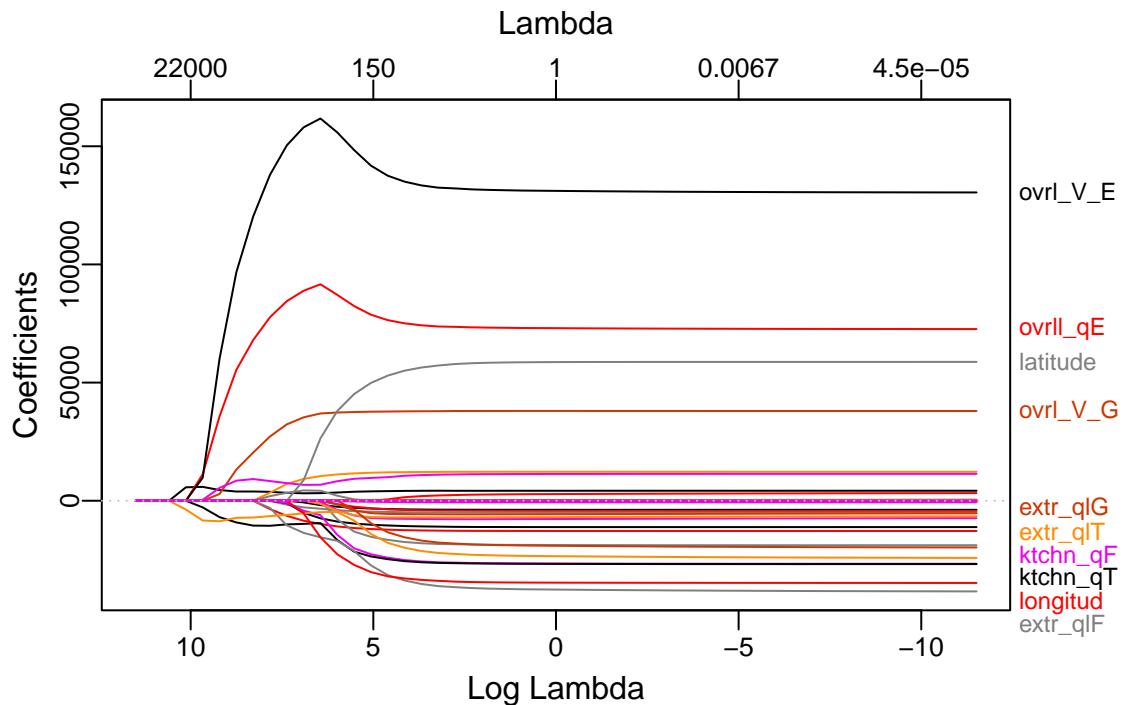
lambda_min = lasso_cv[["lambda.min"]] # minimum mean cross-validated error
lambda_1se = lasso_cv[["lambda.1se"]] # largest value of lambda such that error is within 1 standard er...
```

The λ value with the smallest CVM (5.2673109×10^8) is 39.8107171.

```
plot(lasso_cv)
```



```
plot_glmnet(lasso_cv$glmnet.fit)
```



Test Error with $\lambda = 39.81$:

```
y_pred = predict(lasso_cv, newx = x_test, s = lambda_min, type = "response")
mspe_lasso <- mean((y_test - y_pred)^2)
```

The test error is 4.4304636×10^8 .

```

coef_1se = predict(lasso_cv, type = "coefficients", s = lambda_1se)
# Count non-zero coefficients to determine the number of predictors (excluding intercept)
num_predictors_1se = sum(coef_1se != 0) - 1

```

When using λ_{1SE} , there are 31 predictors.

b - Elastic Net Model

```

set.seed(1234)

alpha = seq(0, 1, length = 21) # alpha grid ranging from 0 (Ridge) to 1 (Lasso)
lambda = 10^(seq(-5, 5, 0.2)) # lambda grid of lambda values for penalty tuning

# Iterate through each alpha to perform cross-validation and store results in a tibble
enet_cv_results = tibble(alpha = alpha) |>
  mutate(
    cv_fit = map(alpha, ~cv.glmnet(x, y, alpha = .x, lambda = lambda)), # runs glmnet for each alpha
    min_cvm = map_dbl(cv_fit, ~min(.x$cvm)) # finds min CVM for each model at each alpha
  )

# pull the alpha value with the lowest overall CVM
enet_alpha_min_cvm = enet_cv_results |>
  filter(min_cvm == min(min_cvm)) |>
  pull(alpha)

# pull the corresponding model with that optimal alpha
best_enet_cv_fit = enet_cv_results |>
  filter(alpha == enet_alpha_min_cvm) |>
  pull(cv_fit) |>
  pluck(1)

# pull 1SE lambda for optimal alpha
lambda_1se_enet = best_enet_cv_fit$lambda.1se

# make predictions using optimal alpha and 1SE lambda at that alpha
y_pred = predict(best_enet_cv_fit, newx = x_test, s = "lambda.1se")
mspe_elastic_net = mean((y_test - y_pred)^2)

```

The selected tuning parameters are $\alpha = 0.45$ and $\lambda_{1SE} = 1584.89$. The model with these parameters has a test error of 4.2029165×10^8 .

Applying the 1SE rule is a bit more complicated with elastic net, because it has two parameters, α and λ . The 1SE rule, when used in the Lasso model, is easy to implement because a larger value for λ means fewer non-zero coefficients and a more parsimonious model. In elastic net models, this will work when α is fixed at a single value.

c - Partial Least Squares (PLS) Model

```

set.seed(1234)

pls_mod <- plsr(sale_price ~ .,
                  data = housing_train,
                  scale = TRUE, # similar scaling importance as PCR

```

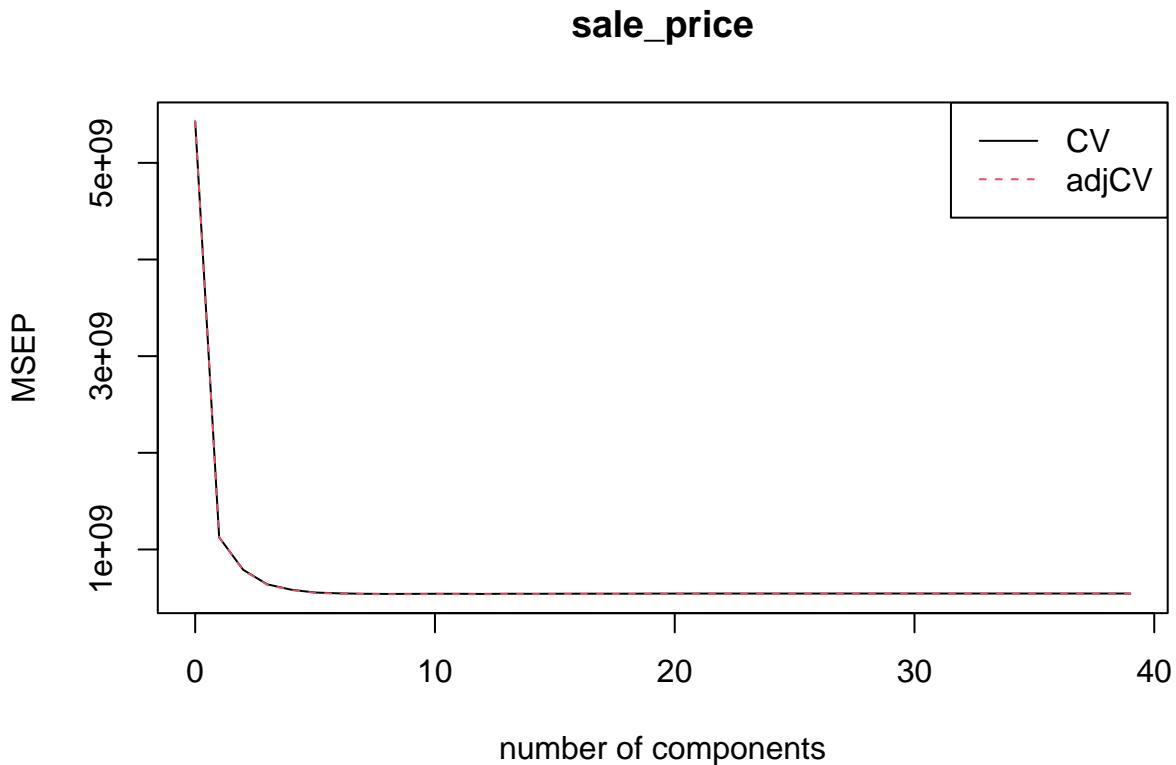
```

    validation = "CV")
summary(pls_mod)

## Data: X dimension: 1440 39
## Y dimension: 1440 1
## Fit method: kernelpls
## Number of components considered: 39
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##          (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV        73685    33553    28106    25289    24162    23546    23362
## adjCV    73685    33537    28060    25207    24086    23471    23295
##          7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV        23277    23238    23250    23272    23269    23240    23282
## adjCV    23210    23173    23182    23200    23196    23170    23207
##          14 comps 15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## CV        23266    23279    23295    23290    23294    23312    23315
## adjCV    23193    23205    23219    23215    23219    23235    23238
##          21 comps 22 comps 23 comps 24 comps 25 comps 26 comps 27 comps
## CV        23323    23322    23322    23322    23323    23324    23326
## adjCV    23245    23245    23244    23244    23245    23246    23248
##          28 comps 29 comps 30 comps 31 comps 32 comps 33 comps 34 comps
## CV        23326    23326    23326    23327    23327    23327    23327
## adjCV    23248    23248    23248    23248    23248    23248    23248
##          35 comps 36 comps 37 comps 38 comps 39 comps
## CV        23327    23327    23327    23327    23326
## adjCV    23248    23248    23248    23248    23266
##
## TRAINING: % variance explained
##          1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps
## X        20.02   25.93   29.67   33.59   37.01   40.03   42.49
## sale_price 79.73   86.35   89.36   90.37   90.87   90.99   91.06
##          8 comps 9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X        45.53   47.97   50.15   52.01   53.69   55.35   56.86
## sale_price 91.08   91.10   91.13   91.15   91.15   91.16   91.16
##          15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## X        58.64   60.01   62.18   63.87   65.26   67.10
## sale_price 91.16   91.16   91.16   91.16   91.16   91.16
##          21 comps 22 comps 23 comps 24 comps 25 comps 26 comps
## X        68.44   70.12   71.72   73.35   75.20   77.27
## sale_price 91.16   91.16   91.16   91.16   91.16   91.16
##          27 comps 28 comps 29 comps 30 comps 31 comps 32 comps
## X        78.97   80.10   81.83   83.55   84.39   86.34
## sale_price 91.16   91.16   91.16   91.16   91.16   91.16
##          33 comps 34 comps 35 comps 36 comps 37 comps 38 comps
## X        88.63   90.79   92.79   95.45   97.49   100.00
## sale_price 91.16   91.16   91.16   91.16   91.16   91.16
##          39 comps
## X        100.24
## sale_price 91.14

# plot cross-validated MSEP for PLS
validationplot(pls_mod, val.type = "MSEP", legendpos = "topright")

```



```
# determine the optimal number of components
cv_mse <- RMSEP(pls_mod)
ncomp_cv <- which.min(cv_mse$val[1, ,]) - 1
ncomp_cv

## 8 comps
##     8

# calculate test MSE
predy2_pls <- predict(pls_mod, newdata = housing_test,
                      ncomp = ncomp_cv)

mspe_pls = mean((y_test - predy2_pls)^2)
```

There are 8 components in the partial least squares model, with an MSPE of 4.4021794×10^8 .

d - Comparing Models

```
summary = tibble(
  model = c("LASSO", "Elastic Net", "PLS"),
  mspe = c(mspe_lasso, mspe_elastic_net, mspe_pls)
)

knitr::kable(summary)
```

model	mspe
LASSO	443046363
Elastic Net	420291645
PLS	440217938

Comparing the mean squared predicted error across the three models, the elastic net model is the best model for making predictions. This means that the model likely benefits from the balance of LASSO and Ridge (λ , α) penalties, as compared to the LASSO model, which only has the λ penalty.