

HW 1

Shivalika Chavan

2026-02-15

```
housing_train = read.csv("./data/housing_training.csv") |> janitor::clean_names()
housing_test = read.csv("./data/housing_test.csv") |> janitor::clean_names()

ordinal_vars = c("Overall_Qual",
                 "Kitchen_Qual",
                 "Fireplace_Qu",
                 "Exter_Qual")

for (v in ordinal_vars) {
  if (v %in% colnames(housing_train)) {
    housing_train[[v]] <- as.ordered(housing_train[[v]])
    housing_test[[v]] <- as.ordered(housing_test[[v]])
  }
}

y = housing_train |> pull(sale_price)
x = model.matrix(sale_price ~ ., housing_train)[,-1]

x_test = model.matrix(sale_price ~ ., housing_test)[,-1]
y_test = housing_test$sale_price
```

a - LASSO Model

```
set.seed(1234)

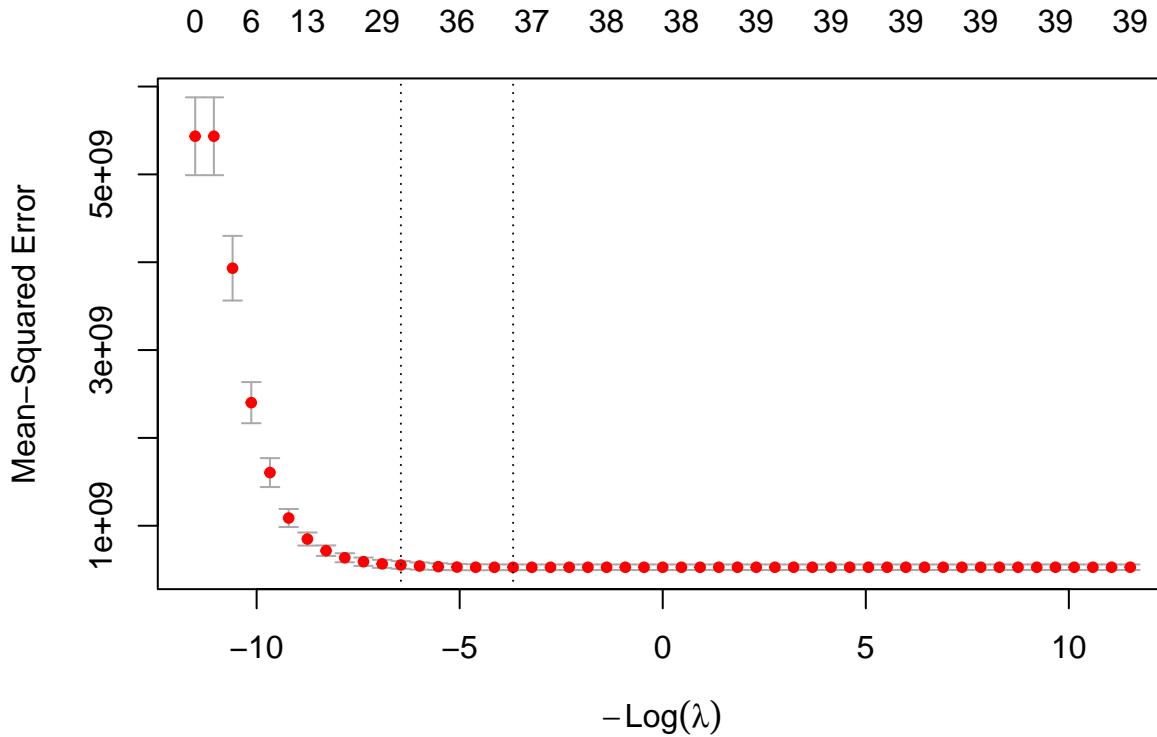
lambda = 10^(seq(-5, 5, 0.2)) # lambda grid of lambda values for penalty tuning

# k-fold cross-validation for Lasso (alpha = 1) over the lambda values
lasso_cv = cv.glmnet(x, y,
                      alpha = 1,
                      lambda = lambda)

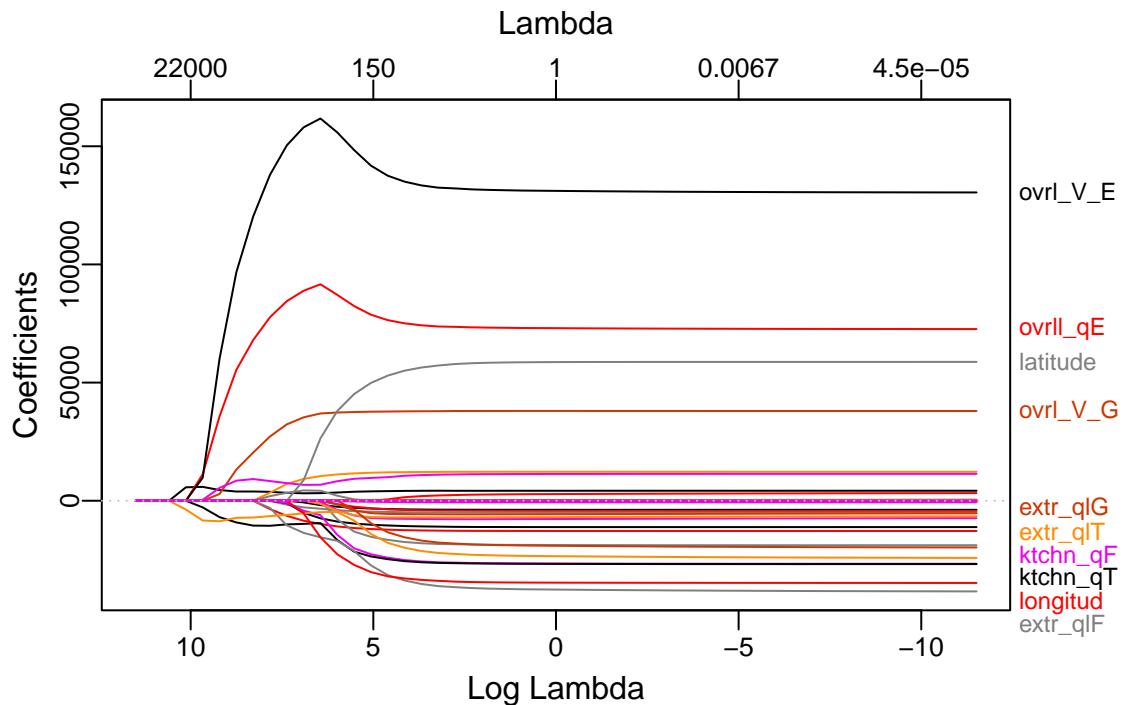
lambda_min_glmnet = lasso_cv[["lambda.min"]] # minimum mean cross-validated error
lambda_1se_glmnet = lasso_cv[["lambda.1se"]] # largest value of lambda such that error is within 1 standard error
CVM_min = lasso_cv |> broom::tidy() |> filter(lambda == lambda_min_glmnet) |> pull(estimate)
CVM_1se = lasso_cv |> broom::tidy() |> filter(lambda == lambda_1se_glmnet) |> pull(estimate)
```

The λ value with the smallest CVM (5.2673109×10^8) is 39.8107171. The λ value with the 1SE CVM (5.5321962×10^8) is 630.9573445.

```
plot(lasso_cv)
```



```
plot_glmnet(lasso_cv$glmnet.fit)
```



Test Error with $\lambda = 39.81$:

```
y_pred = predict(lasso_cv, newx = x_test, s = lambda_min_glmnet, type = "response")
mspe_lasso <- mean((y_test - y_pred)^2)
```

The test error is 4.4304636×10^8 .

```

coef_1se = predict(lasso_cv, type = "coefficients", s = lambda_1se_glmnet)
# Count non-zero coefficients to determine the number of predictors (excluding intercept)
num_predictors_1se = sum(coef_1se != 0) - 1

```

When using λ_{1SE} , there are 31 predictors.

b - Elastic Net Model

```

set.seed(1234)

alpha = seq(0, 1, length = 21) # alpha grid ranging from 0 (Ridge) to 1 (Lasso)
lambda = 10^(seq(-5, 5, 0.2)) # lambda grid of lambda values for penalty tuning

# Iterate through each alpha to perform cross-validation and store results in a tibble
enet_cv_results = tibble(alpha = alpha) |>
  mutate(
    cv_fit = map(alpha, ~cv.glmnet(x, y, alpha = .x, lambda = lambda)), # runs glmnet for each alpha
    min_cvm = map_dbl(cv_fit, ~min(.x$cvm)) # finds min CVM for each model at each alpha
  )

# pull the alpha value with the lowest overall CVM
enet_alpha_min_cvm = enet_cv_results |>
  filter(min_cvm == min(min_cvm)) |>
  pull(alpha)

# pull the corresponding model with that optimal alpha
best_enet_cv_fit = enet_cv_results |>
  filter(alpha == enet_alpha_min_cvm) |>
  pull(cv_fit) |>
  pluck(1)

# pull 1SE lambda for optimal alpha
lambda_1se_enet = best_enet_cv_fit$lambda.1se

# make predictions using optimal alpha and 1SE lambda at that alpha
y_pred = predict(best_enet_cv_fit, newx = x_test, s = "lambda.1se")
mspe_elastic_net = mean((y_test - y_pred)^2)

```

The selected tuning parameters are $\alpha = 0.45$ and $\lambda_{1SE} = 1584.89$. The model with these parameters has a test error of 4.2029165×10^8 .

Applying the 1SE rule is a bit more complicated with elastic net, because it has two parameters, α and λ . The 1SE rule, when used in the Lasso model, is easy to implement because a larger value for λ means fewer non-zero coefficients and a more parsimonious model. In elastic net models, this will work when α is fixed at a single value.

c - Partial Least Squares (PLS) Model

```

set.seed(1234)

pls_mod <- plsr(sale_price ~ .,
                  data = housing_train,
                  scale = TRUE, # similar scaling importance as PCR

```

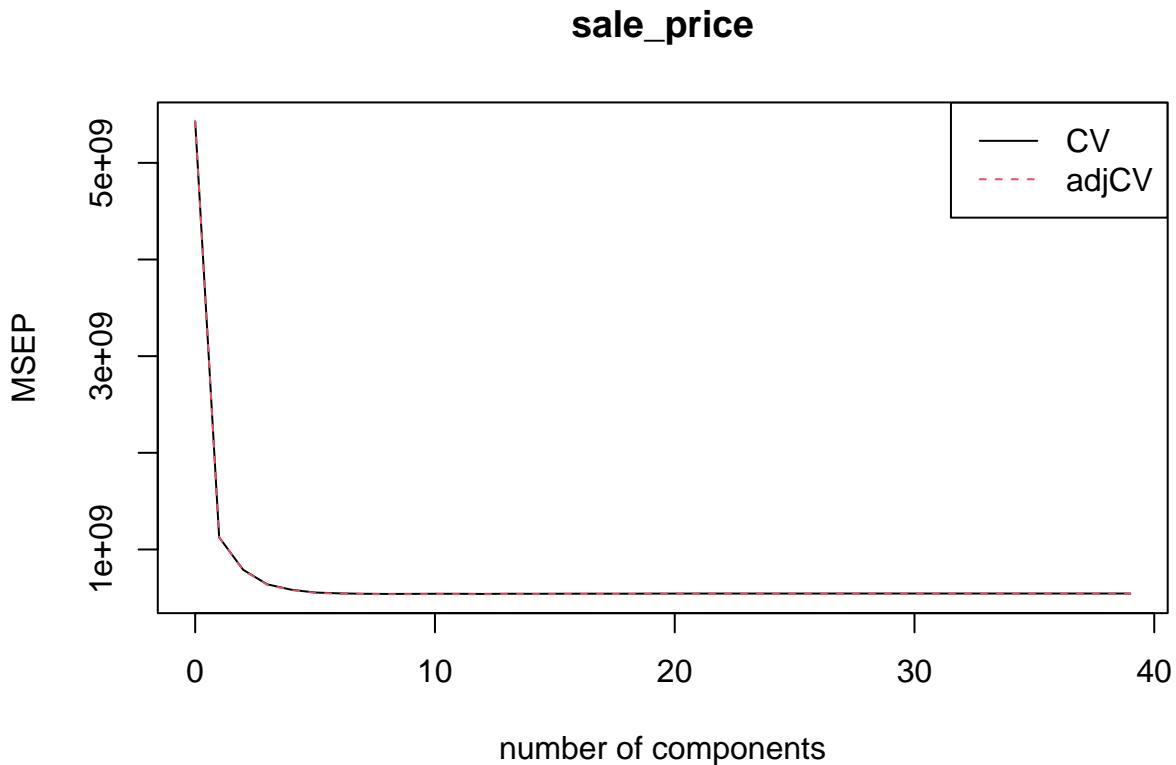
```

        validation = "CV")
summary(pls_mod)

## Data:      X dimension: 1440 39
## Y dimension: 1440 1
## Fit method: kernelpls
## Number of components considered: 39
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##          (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV          73685   33553   28106   25289   24162   23546   23362
## adjCV       73685   33537   28060   25207   24086   23471   23295
##          7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV          23277   23238   23250   23272   23269   23240   23282
## adjCV       23210   23173   23182   23200   23196   23170   23207
##          14 comps 15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## CV          23266   23279   23295   23290   23294   23312   23315
## adjCV       23193   23205   23219   23215   23219   23235   23238
##          21 comps 22 comps 23 comps 24 comps 25 comps 26 comps 27 comps
## CV          23323   23322   23322   23322   23323   23324   23326
## adjCV       23245   23245   23244   23244   23245   23246   23248
##          28 comps 29 comps 30 comps 31 comps 32 comps 33 comps 34 comps
## CV          23326   23326   23326   23327   23327   23327   23327
## adjCV       23248   23248   23248   23248   23248   23248   23248
##          35 comps 36 comps 37 comps 38 comps 39 comps
## CV          23327   23327   23327   23327   23326
## adjCV       23248   23248   23248   23248   23266
##
## TRAINING: % variance explained
##          1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps
## X          20.02   25.93   29.67   33.59   37.01   40.03   42.49
## sale_price 79.73   86.35   89.36   90.37   90.87   90.99   91.06
##          8 comps 9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X          45.53   47.97   50.15   52.01   53.69   55.35   56.86
## sale_price 91.08   91.10   91.13   91.15   91.15   91.16   91.16
##          15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## X          58.64   60.01   62.18   63.87   65.26   67.10
## sale_price 91.16   91.16   91.16   91.16   91.16   91.16
##          21 comps 22 comps 23 comps 24 comps 25 comps 26 comps
## X          68.44   70.12   71.72   73.35   75.20   77.27
## sale_price 91.16   91.16   91.16   91.16   91.16   91.16
##          27 comps 28 comps 29 comps 30 comps 31 comps 32 comps
## X          78.97   80.10   81.83   83.55   84.39   86.34
## sale_price 91.16   91.16   91.16   91.16   91.16   91.16
##          33 comps 34 comps 35 comps 36 comps 37 comps 38 comps
## X          88.63   90.79   92.79   95.45   97.49   100.00
## sale_price 91.16   91.16   91.16   91.16   91.16   91.16
##          39 comps
## X          100.24
## sale_price 91.14

# plot cross-validated MSEP for PLS
validationplot(pls_mod, val.type = "MSEP", legendpos = "topright")

```



```
# determine the optimal number of components
cv_mse <- RMSEP(pls_mod)
ncomp_cv <- which.min(cv_mse$val[1, ,]) - 1
ncomp_cv

## 8 comps
##     8

# calculate test MSE
predy2_pls <- predict(pls_mod, newdata = housing_test,
                      ncomp = ncomp_cv)

mspe_pls = mean((y_test - predy2_pls)^2)
```

There are 8 components in the partial least squares model, with an MSPE of 4.4021794×10^8 .

d - Comparing Models

```
summary = tibble(
  model = c("LASSO", "Elastic Net", "PLS"),
  mspe = c(mspe_lasso, mspe_elastic_net, mspe_pls)
)

knitr::kable(summary)
```

model	mspe
LASSO	443046363
Elastic Net	420291645
PLS	440217938

Comparing the mean squared predicted error across the three models, the elastic net model is the best model for making predictions. This means that the model likely benefits from the balance of LASSO and Ridge (λ , α) penalties, as compared to the LASSO model, which only has the λ penalty.

e - Retraining LASSO model using caret instead of glmnet

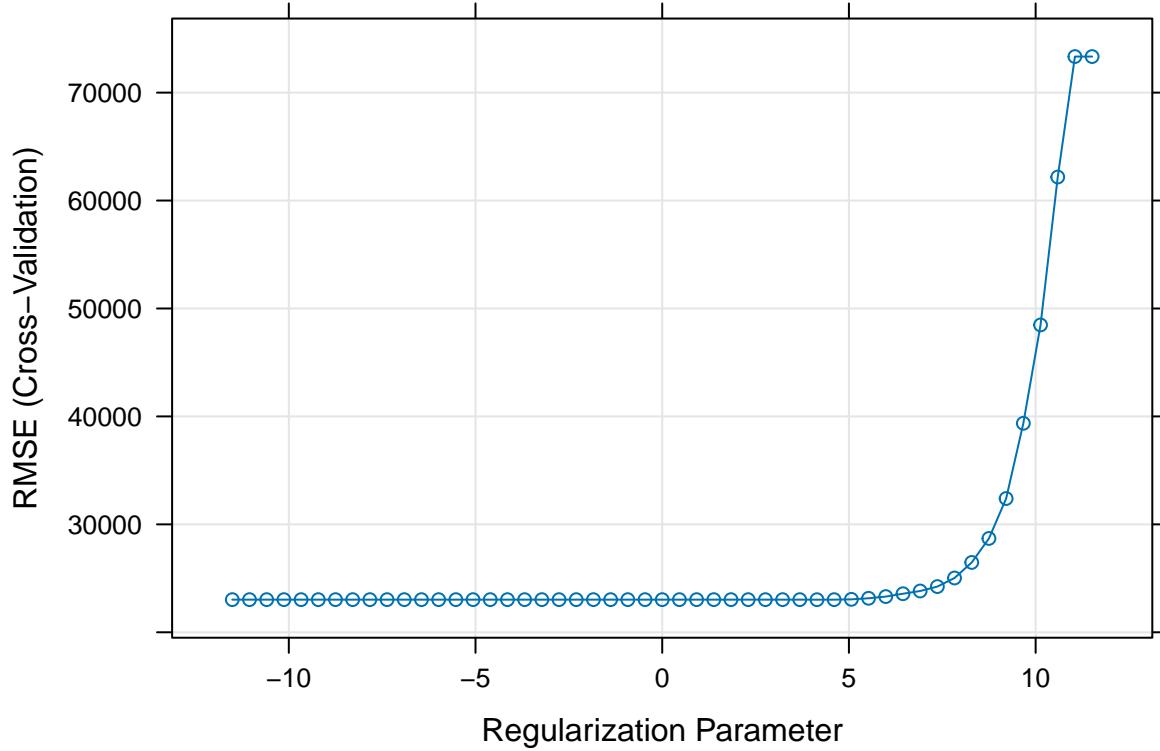
```
set.seed(1234)

ctrl1 = trainControl(method = "cv", number = 10) #cross validation, in this case 10-fold

lasso_caret = train(sale_price ~ .,
                    data = housing_train,
                    method = "glmnet",
                    tuneGrid = expand.grid(alpha = 1,
                                           lambda = lambda),
                    trControl = ctrl1)

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.

plot(lasso_caret, xTrans = log)
```



```
lambda_min_caret = lasso_caret$bestTune$lambda

tibble(
  package = c("caret (min)", "glmnet (min)", "glmnet (1SE)"),
  lambda = c(lambda_min_caret, lambda_min_glmnet, lambda_1se_glmnet)
) |> knitr::kable(digits = 2)
```

package	lambda
caret (min)	63.10
glmnet (min)	39.81
glmnet (1SE)	630.96

The “best” values for λ differ slightly between `caret` and `cv.glmnet`. However, these two values are much closer to each other than they are to the 1SE value found using `cv.glmnet`. While they will result in slightly different models, their cross-validation errors will be within 1SE of each other. The reason for this discrepancy is the difference in internal functions; specifically, how each package assigns observations to k cross-validation folds, even when the random seed `set.seed(1234)` is kept consistent.